

# Lecture: Logistic Regression

## Binary classification, sigmoid and model tuning

K. Kadri

# Outline

- 1 Intuition and role of logistic regression
- 2 Mathematical formulation
- 3 Regularization and hyperparameters
- 4 Decision threshold and metrics
- 5 Advantages, limitations and best practices

## Problem: predicting a binary class

- We want to predict a variable  $y \in \{0, 1\}$  (e.g., *sick / not sick*).
- A linear regression would output an unbounded real value.
- We need a model that outputs a **probability** between 0 and 1:

$$p(y = 1 | x) \in [0, 1]$$

### Key idea

Logistic regression does not directly predict a class, but the **probability of the positive class**, then applies a **threshold** (often 0.5).

# The sigmoid function

## Definition

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

where  $z = w^\top x + b$ .

- For  $z \rightarrow +\infty : \sigma(z) \rightarrow 1$ .
- For  $z \rightarrow -\infty : \sigma(z) \rightarrow 0$ .
- Interpretation:  $\sigma(z)$  is a **probability**.

## Logistic model

$$p(y = 1 | x) = \sigma(w^\top x + b)$$

# Log-odds and linear model

- The model can be rewritten as **log-odds**:

$$\log \frac{p}{1-p} = \mathbf{w}^\top \mathbf{x} + b$$

- $\frac{p}{1-p}$ : **odds ratio**.
- The algorithm learns  $w$  and  $b$  that explain this odds ratio.

## Interpretation

Each coefficient  $w_j$  measures the effect of feature  $x_j$  on the **log-odds** of belonging to the positive class.

# Cost function: log-loss

For one example  $(x_i, y_i)$

$$\mathcal{L}(y_i, \hat{p}_i) = -[y_i \log(\hat{p}_i) + (1 - y_i) \log(1 - \hat{p}_i)]$$

where  $\hat{p}_i = p(y = 1 | x_i)$ .

- If the model is confident and correct  $\rightarrow$  low loss.
- If the model is confident and wrong  $\rightarrow$  very high loss.

## Learning objective

Minimize the sum (or mean) of the log-loss over the entire training set.

# L2 and L1 regularization

- **Problem:** overfitting if coefficients have too much freedom.
- **L2 regularization (Ridge):**

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{log-loss}} + \lambda \|\mathbf{w}\|_2^2$$

- **L1 regularization (Lasso):**

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{log-loss}} + \lambda \|\mathbf{w}\|_1$$

## Effect

- L2: shrinks all coefficients (rarely to 0).
- L1: pushes some coefficients exactly to 0 → **feature selection.**

# Hyperparameter $C$ in scikit-learn

## In LogisticRegression

$$C = \frac{1}{\lambda}$$

- Small  $C \rightarrow$  strong regularization (simpler, more biased model).
- Large  $C \rightarrow$  weak regularization (more complex, higher overfitting risk).

## Tuning

We choose  $C$  (and penalty type) via **cross-validation** (GridSearch, RandomizedSearch).

# Decision threshold

- Default:

$$\hat{y} = \begin{cases} 1 & \text{if } p(y = 1 | x) \geq 0.5 \\ 0 & \text{otherwise} \end{cases}$$

- But we can **adapt the threshold** depending on the problem:

- lower threshold → more positives detected (recall ↑, FP risk ↑).
- higher threshold → fewer false positives but lower recall.

## Link with metrics

We choose the threshold based on the **trade-off** Precision / Recall / F1-score and the **business cost** of false positives and false negatives.

# ROC curve and AUC

- By varying the threshold, we obtain multiple (FPR, TPR) points.
- The ROC curve plots:  
TPR as a function of FPR
- The area under the curve (AUC) measures the model's **overall ability** to separate classes.

## Rule of thumb

- $AUC \approx 0.5$ : random model.
- $AUC$  close to 1: very good separation.

# Advantages of logistic regression

- **Simple** model, fast to train.
- **Probabilistic** output, easy to interpret.
- Interpretable: coefficients can be analyzed.
- Works well as a **baseline** on many problems.

## Typical use cases

- Credit scoring.
- Customer churn probability.
- Binary diagnosis (presence/absence of a disease).

- Assumes a **linear decision boundary** in feature space.
- Sensitive to **poorly scaled** features → standardization important.
- May be insufficient if the relationship is highly nonlinear.

## Best practices

- Always standardize continuous features.
- Test several  $C$  values and penalty types.
- Examine coefficients and metrics (F1, AUC) rather than accuracy alone.