

# Machine Learning Algorithm: Naive Bayes Classifier

## Bayes' theorem, independence assumption, concrete examples

K. Kadri

# Outline

- 1 Intuition and Motivation
- 2 Naive Independence Assumption
- 3 Concrete Example: Spam Classification
- 4 Training and Prediction Algorithm
- 5 Laplace Smoothing
- 6 Advantages, Limits, Best Practices
- 7 Visual Schema and Variants

# Problem: predicting a class

- We want to predict a label  $y \in \{0, 1\}$  or multiple classes.
- Example:
  - Email: **spam** / **not spam**
  - Medical test: **disease** / **no disease**
- We need a model that uses available features to estimate:

$$p(y | x)$$

## Key idea

Naive Bayes predicts the class with the highest **posterior probability** based on Bayes' theorem.

# Bayes' Theorem

## Definition

$$p(y | x) = \frac{p(x | y) p(y)}{p(x)}$$

- $p(y)$  : prior probability of the class
- $p(x | y)$  : likelihood
- $p(y | x)$  : posterior probability

## Goal

Choose the class with the highest posterior:

$$\hat{y} = \arg \max_y p(y | x)$$

# Naive assumption

- Exact computation of  $p(x | y)$  is hard:

$$x = (x_1, x_2, \dots, x_n)$$

- Naive Bayes assumes:

$$p(x | y) = \prod_{j=1}^n p(x_j | y)$$

## Meaning

Features are assumed **conditionally independent** given the class.

## Resulting classifier

$$p(y | x) \propto p(y) \prod_{j=1}^n p(x_j | y)$$

## Example: spam vs. not spam

Suppose we observe the words:

“free”, “money”

- We estimate from training emails:

- $p(\text{spam}) = 0.3$
- $p(\text{not spam}) = 0.7$
- $p(\text{"free"} \mid \text{spam}) = 0.8$
- $p(\text{"money"} \mid \text{spam}) = 0.6$
- $p(\text{"free"} \mid \text{not spam}) = 0.1$
- $p(\text{"money"} \mid \text{not spam}) = 0.05$

$$p(\text{spam} \mid x) \propto 0.3 \times 0.8 \times 0.6 = 0.144$$

$$p(\text{not spam} \mid x) \propto 0.7 \times 0.1 \times 0.05 = 0.0035$$

### Prediction

Email classified as **spam**.

# Training Naive Bayes

## Training step

- Estimate class prior:

$$p(y) = \frac{\text{count}(y)}{\text{total samples}}$$

- Estimate likelihoods:

$$p(x_j | y) = \frac{\text{count}(x_j, y)}{\text{count}(y)}$$

## Efficiency

Training is extremely fast:

- no gradient descent
- no matrix inversion
- scalable to large datasets

# Prediction Algorithm

For a new sample  $x$

Compute:

$$p(y) \prod_{j=1}^n p(x_j | y)$$

$$\hat{y} = \arg \max_y p(y) \prod_{j=1}^n p(x_j | y)$$

## Decision rule

Choose the class with the highest posterior probability.

## Problem: zero probability

- If a word never appears in spam,

$$p(\text{word} \mid \text{spam}) = 0$$

- Then:

$$p(\text{spam} \mid x) = 0$$

- One word can ruin classification.

## Solution: Laplace smoothing

$$p(x_j \mid y) = \frac{\text{count}(x_j, y) + 1}{\text{count}(y) + K}$$

where  $K$  = number of possible values.

# Advantages

- Extremely fast training
- Works well with text data
- Probabilistic output
- Robust with few samples

## Typical use cases

- Spam detection
- Sentiment analysis
- Medical diagnosis
- Document classification

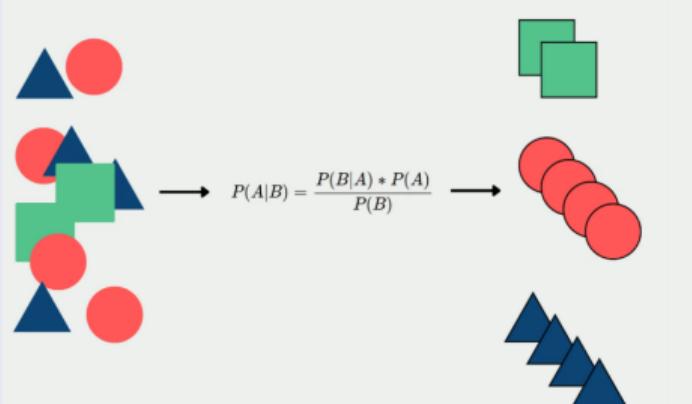
- Independence assumption often unrealistic
- Performs poorly with correlated features
- Continuous features require distributions (Gaussian NB)

## Best practices

- Use TF-IDF for text
- Apply Laplace smoothing
- Compare with logistic regression / SVM

# From Features → Bayes Theorem → Prediction

## Flow of the Naive Bayes Classifier



## Interpretation

The classifier transforms raw features into class probabilities using Bayes' theorem, then selects the most likely class.

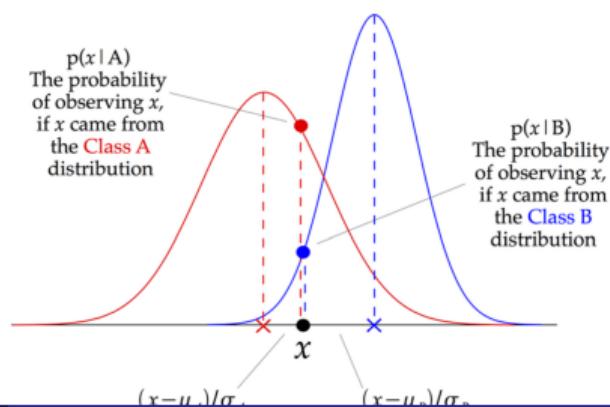
# Gaussian Naive Bayes: Intuition

## Idea

For continuous features, we assume:

$$x_j | y \sim \mathcal{N}(\mu_{y,j}, \sigma_{y,j}^2)$$

- Each class models features as bell-shaped Gaussian curves
- Classes differ by mean and variance
- Probability estimated from the curve height



# Naive Bayes vs Logistic Regression

## Naive Bayes

- Generative model
- Learns:  $p(x | y)$  and  $p(y)$
- Fast, works well with text
- Strong independence assumption

## Logistic Regression

- Discriminative model
- Learns:  $p(y | x)$  directly
- More flexible
- Needs optimization and feature scaling

# Naive Bayes vs Logistic Regression

## Rule of thumb

Naive Bayes performs surprisingly well with sparse text features, while Logistic Regression usually wins when independence is violated.